

Andrey Ustyuzhanin



PyTorch.Next

Auto differentiation, data handling and code disentanglement

2021



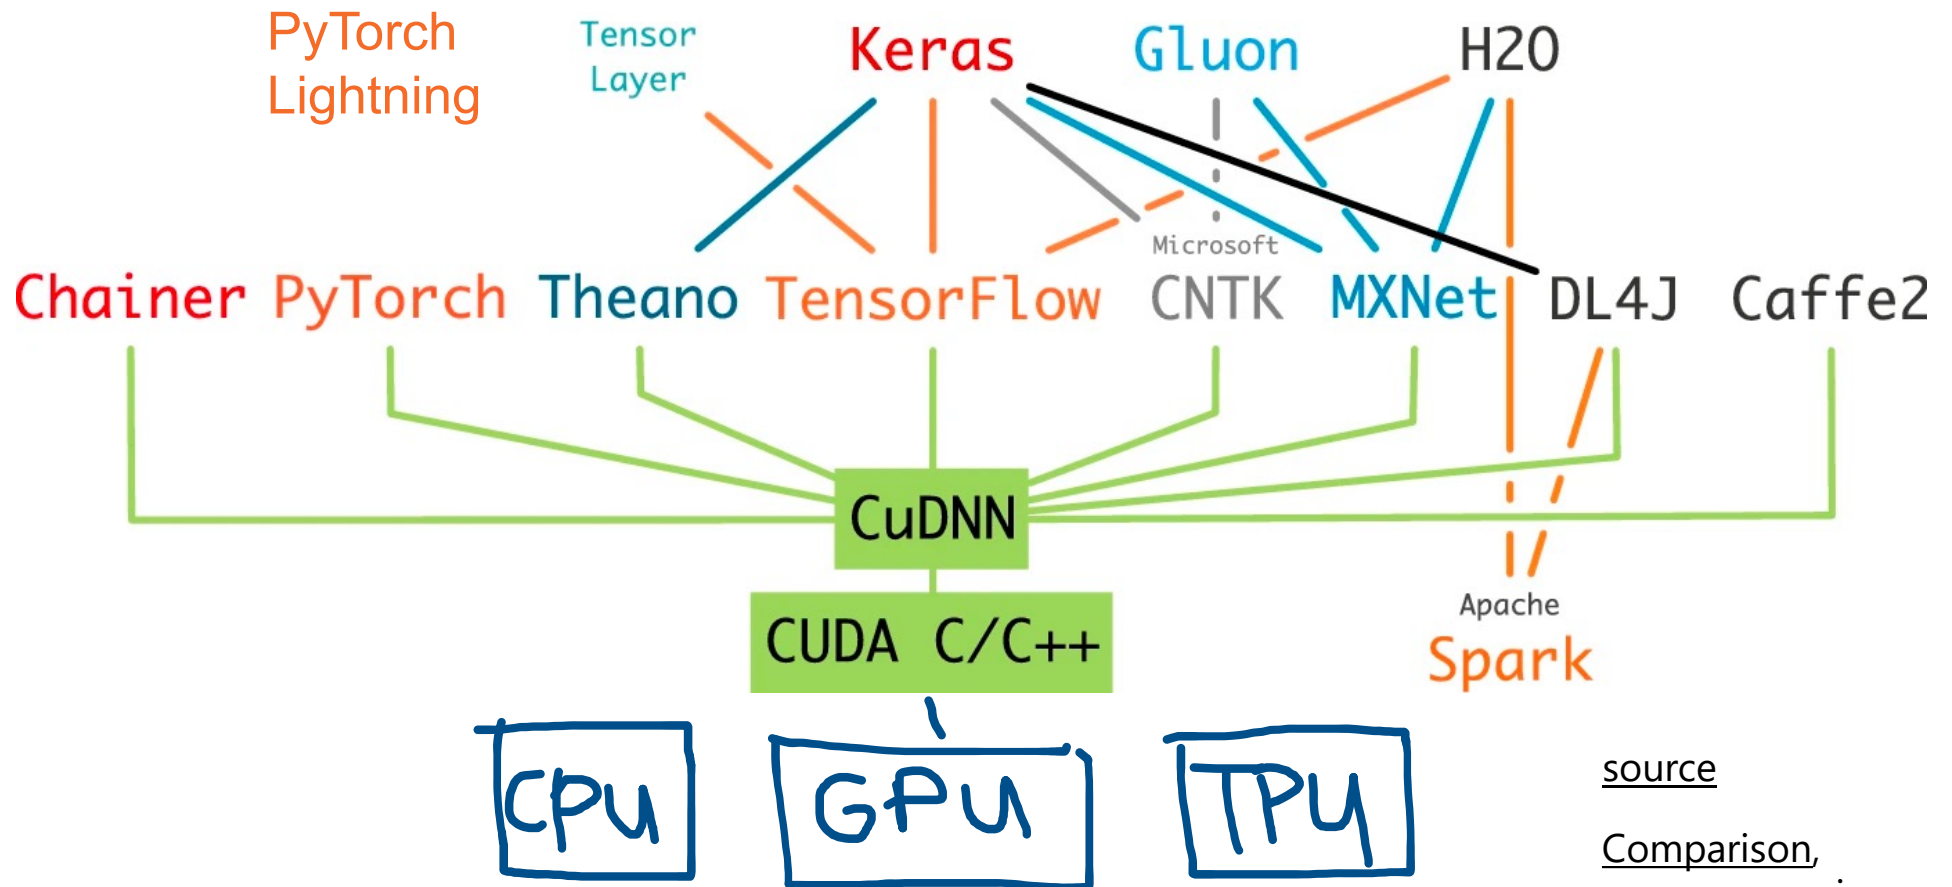
Yandex



EPFL

S_{IT}
Schaffhausen
Institute of
Technology

Deep Learning Frameworks




source

Comparison,
moar comparison

OPINIONS DIFFER
AMONG THE
SAINTS...

AS WELL AS AMONG
THE SINNERS.

-ABRAHAM
LINCOLN

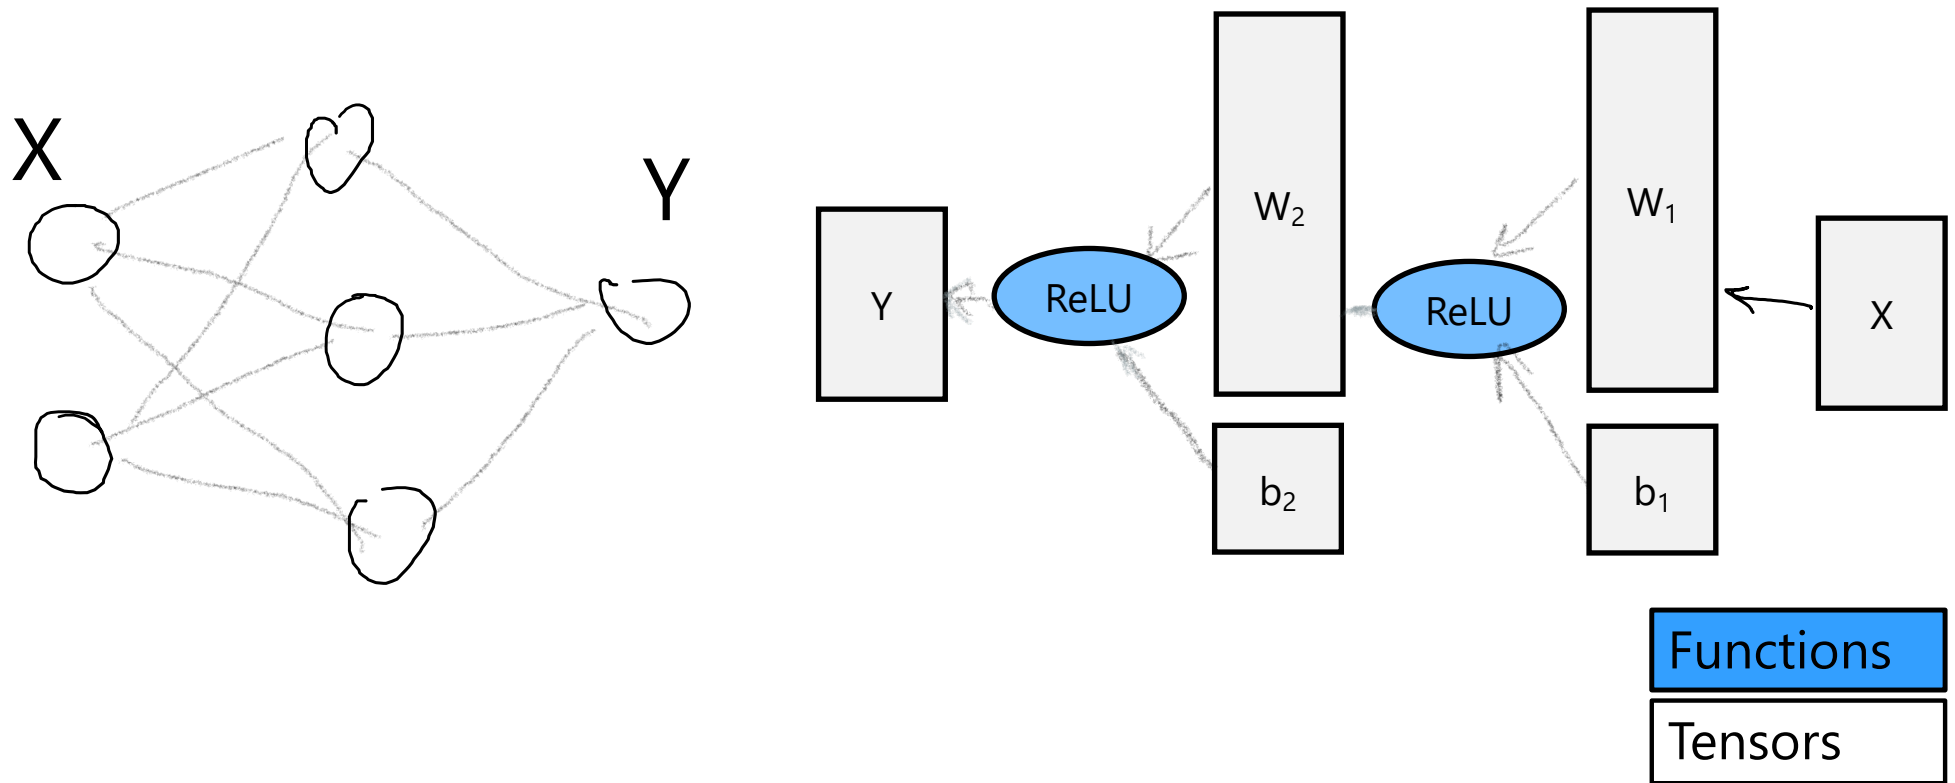


PyTorch highlights

- ▶ Simple, transparent development/ debugging
- ▶ Rich Ecosystem:
 - Plenty of pretrained models
 - NLP, Vision, ...
 - Interpretation
 - Hyper-optimization
- ▶ Production Ready (C++, ONNX, Services)
- ▶ Distributed Training, declarative data parallelism
- ▶ Cloud Deployment support
- ▶ Choice of many industry leaders and researchers

facebook Artificial Intelligence

Neural network representation



$$Y = \text{relu}(W_2 \times \text{relu}(W_1 X + b_1) + b_2)$$

Building blocks, tensors

<code>torch.randn(*size)</code>	<i># tensor with independent $N(0,1)$ entries</i>
<code>torch.ones zeros>(*size)</code>	<i># tensor with all 1's [or 0's]</i>
<code>torch.Tensor(L)</code>	<i># create tensor from [nested] list or ndarray L</i>
<code>x.clone()</code>	<i># clone of x</i>
<code>with torch.no_grad():</code>	<i># code wrap that stops autograd from tracking tensor history</i>
<code>requires_grad=True</code>	<i># arg, when set to True, tracks computation</i>
	<i># history for future derivative calculations</i>
<code>x.size()</code>	<i># return tuple-like object of dimensions</i>
<code>torch.cat(tensor_seq, dim=0)</code>	<i># concatenates tensors along dim</i>
<code>x.view(a,b,...)</code>	<i># reshapes x into size (a,b,...)</i>
<code>x.view(-1,a)</code>	<i># reshapes x into size (b,a) for some b</i>
<code>x.transpose(a,b)</code>	<i># swaps dimensions a and b</i>
<code>x.permute(*dims)</code>	<i># permutes dimensions</i>
<code>x.unsqueeze(dim)</code>	<i># tensor with added axis</i>
<code>x.unsqueeze(dim=2)</code>	<i># (a,b,c) tensor -> (a,b,1,c) tensor</i>

Tensor creation and placement

```
>>> torch.zeros([2, 4], dtype=torch.int32)
tensor([[ 0,  0,  0,  0],
        [ 0,  0,  0,  0]], dtype=torch.int32)
>>> cuda0 = torch.device('cuda:0')
>>> torch.ones([2, 4], dtype=torch.float64, device=cuda0)
tensor([[ 1.0000,  1.0000,  1.0000,  1.0000],
        [ 1.0000,  1.0000,  1.0000,  1.0000]], dtype=torch.float64,
device='cuda:0')
```

- ▶ Keep in mind occurrence of tensors on devices: CPU, GPU, TPU
- ▶ Operations can be performed only if its arguments are inhabiting the same device

GPU, TPU support

```
torch.cuda.is_available          # check for cuda
x.cuda()                        # move x's data from
                                # CPU to GPU and return new object

x.cpu()                          # move x's data from GPU to CPU
                                # and return new object

if not args.disable_cuda and torch.cuda.is_available(): # device agnostic code
    args.device = torch.device('cuda')                # and modularity
else:                                                  #
    args.device = torch.device('cpu')                  #

net.to(device)                                         # recursively convert their
                                                        # parameters and buffers to
                                                        # device specific tensors

mytensor.to(device)                                   # copy your tensors to a device
                                                        # (gpu, cpu)
```

- ▶ <https://pytorch.org/docs/stable/cuda.html>
- ▶ <http://pytorch.org/xla/release/1.5/index.html>

Building blocks, graph


Toy example:

$$b = w_1 * a$$

$$c = w_2 * a$$

$$d = w_3 * b + w_4 * c$$

$$L = 10 - d$$



source

Math operations

```
A.mm(B)      # matrix multiplication  
A.mv(x)      # matrix-vector multiplication  
x.t()        # matrix transpose
```

- ▶ <https://pytorch.org/docs/stable/torch.html?highlight=mm#math-operations>

Computing backpropagation




$$b = w_1 * a$$

$$c = w_2 * a$$

$$d = w_3 * b + w_4 * c$$

$$L = 10 - d$$

$$\frac{\partial L}{\partial w_1} = \frac{\partial L}{\partial d} * \frac{\partial d}{\partial b} * \frac{\partial b}{\partial w_1}$$



Computing gradient automatically

```
>> t1 = torch.randn((3,3), requires_grad = True)

>> t2 = torch.FloatTensor(3,3) # No way to specify requires_grad
>> t2.requires_grad = True
```

Each **Tensor** has an attribute **grad_fn**, which refers to the mathematical operator that created it.

If **Tensor** is a leaf node (initialized by the user), then the **grad_fn** is **None**.

```
import torch

a = torch.randn((3,3), requires_grad = True)

w1 = torch.randn((3,3), requires_grad = True)
w2 = torch.randn((3,3), requires_grad = True)
w3 = torch.randn((3,3), requires_grad = True)
w4 = torch.randn((3,3), requires_grad = True)

b = w1*a
c = w2*a

d = w3*b + w4*c

L = 10 - d


print("The grad fn for a is", a.grad_fn)
print("The grad fn for d is", d.grad_fn)
```

```
The grad fn for a is None
The grad fn for d is <AddBackward0 object at 0x1033afe48>
```

Functions

- ▶ All math steps represented by classes inherited from *torch.autograd.Function*

- *forward*, computes node output and buffers it
- *backward*, stores incoming gradient in **grad** and passes further



```
def backward (incoming_gradients):  
    self.Tensor.grad = incoming_gradients  
  
    for inp in self.inputs:  
        if inp.grad_fn is not None:  
            new_incoming_gradients = //  
                incoming_gradient * local_grad(self.Tensor, inp)  
  
            inp.grad_fn.backward(new_incoming_gradients)  
        else:  
            pass
```

Gradient descent

- Compute gradient for every tensor involved

```
import torch

a = torch.randn((3,3), requires_grad = True)

w1 = torch.randn((3,3), requires_grad = True)
w2 = torch.randn((3,3), requires_grad = True)
w3 = torch.randn((3,3), requires_grad = True)
w4 = torch.randn((3,3), requires_grad = True)

b = w1*a
c = w2*a

d = w3*b + w4*c

# Replace L = (10 - d) by
L = (10 - d).sum()

L.backward()
```

- Make gradient descent step in the opposite direction:

```
learning_rate = 0.5
w1 = w1 - learning_rate * w1.grad
```

Dynamic graph

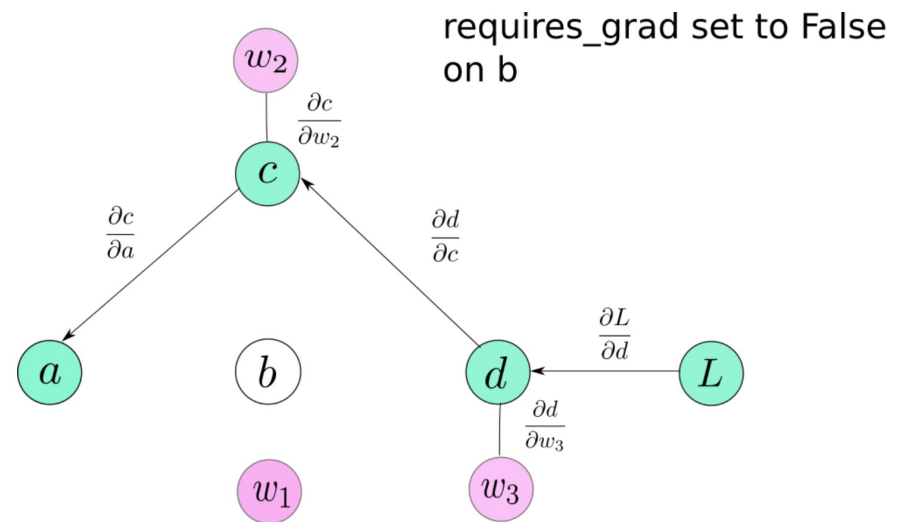
- ▶ Calling *forward* creates
 - graph with the intermediate node output values,
 - buffers for the non-leaf nodes,
 - buffers for intermediate gradient values.
- ▶ Calling *backward*
 - computes gradients and
 - frees the buffers and destroys the graph.
- ▶ Next time, calling *forward*
 - leaf node buffers from the previous run will be shared,
 - **non-leaf nodes buffers will be recreated.**

Gradient cleaning

- ▶ Due to the flexibility of the network architecture, it is not obvious when does iteration of a gradient descent stops, so *backward's* gradients are accumulated each time a variable (Tensor) occurs in the graph;
- ▶ It is usually desired for RNN cases;
- ▶ If you do not need to accumulate those, you must **clean previous gradient values** at the end of each iteration:
 - Either by `x.data.zero_()` for every model tensor `x`;
 - Or by optimizers's `zero_grad()` method, which is more preferable.

Freezing weights

- ▶ **Requires_grad** attribute of the *Tensor* class. By default, it's **False**. It comes handy when you must freeze some layers and stop them from updating parameters while training.
- ▶ Thus, no gradient would be propagated to them, or to those layers which depend upon these layers for gradient flow **requires_grad**.
- ▶ When set to **True**, **requires_grad** is contagious: even if one operand of an operation has **requires_grad** set to **True**, so will the result.



Pre-trained models' enhancement

```
model = torchvision.models.resnet18(pretrained=True)
for param in model.parameters():
    param.requires_grad = False
# Replace the last fully-connected layer
# Parameters of newly constructed modules have requires_grad=True by default
model.fc = nn.Linear(512, 100)

# Optimize only the classifier
optimizer = optim.SGD(model.fc.parameters(), lr=1e-2, momentum=0.9)
```

Inference

- ▶ When we are computing gradients, we need to cache input values, and intermediate features as they may be required to compute the gradient later. The gradient of $\mathbf{b} = \mathbf{w1} * \mathbf{a}$ w.r.t its inputs $\mathbf{w1}$ and \mathbf{a} is \mathbf{a} and $\mathbf{w1}$, respectively.
- ▶ We need to store these values for gradient computation during the backward pass. This affects the memory footprint of the network.
- ▶ While, we are performing inference, we don't compute gradients

```
with torch.no_grad:  
    inference code goes here
```

- ▶ Even better and recent optimized context: `with torch.inference_mode` ([link](#))

Neural Network class: torch.nn.Module

```
import torch.nn as nn
import torch.nn.functional as F

class Model(nn.Module):
    def __init__(self):
        super(Model, self).__init__()
        self.conv1 = nn.Conv2d(1, 20, 5)
        self.conv2 = nn.Conv2d(20, 20, 5)

    def forward(self, x):
        x = F.relu(self.conv1(x))
        return F.relu(self.conv2(x))
```

Loss functions

```
nn.X                                     # where X is BCELoss, CrossEntropyLoss,  
                                         # L1Loss, MSELoss, NLLLoss, SoftMarginLoss,  
                                         # MultiLabelSoftMarginLoss, CosineEmbeddingLoss,  
                                         # KLDivLoss, MarginRankingLoss, HingeEmbeddingLoss  
                                         # or CosineEmbeddingLoss
```

- ▶ <https://pytorch.org/docs/stable/nn.html#loss-functions>

Optimizers

```
opt = optim.x(model.parameters(), ...)    # create optimizer  
opt.step()                               # update weights  
optim.X                                  # where X is SGD, Adadelata, Adagrad, Adam,  
                                           # SparseAdam, Adamax, ASGD,  
                                           # LBFGS, RMSProp or Rprop
```

- ▶ <https://pytorch.org/docs/stable/optim.html>

Data Utils

Datasets

```
Dataset           # abstract class representing dataset  
TensorDataset     # labelled dataset in the form of tensors  
Concat Dataset    # concatenation of Datasets
```

- ▶ <https://pytorch.org/docs/stable/data.html?highlight=dataset#torch.utils.data.Dataset>

Dataloaders and DataSamplers

```
DataLoader(dataset, batch_size=1, ...) # loads data batches agnostic  
                                         # of structure of individual data points  
  
sampler.Sampler(dataset,...)           # abstract class dealing with  
                                         # ways to sample from dataset  
  
sampler.XSampler where ...             # Sequential, Random, Subset,  
                                         # WeightedRandom or Distributed
```

- ▶ <https://pytorch.org/docs/stable/data.html?highlight=dataloader#torch.utils.data.DataLoader>

PyTorch Lightning

- ▶ The lightweight PyTorch wrapper for high-performance AI research.


Maximal flexibility

```
def training_step(self, batch, batch_nb):  
    x, y = batch  
    z = self.encoder(x)  
    x_hat = self.decoder(z)  
    mse = F.mse_loss(x_hat, x)  
    gan_regularizer = self.discriminator(x_hat)  
    loss = mse + gan_regularizer  
    return loss
```

No boilerplate Maximal flexibility

```
if gpu:  
    x = x.cuda(0)  
    z = encoder(x)  
    x_hat = decoder(z)  
    backward()
```

Self contained models



Modular



Research
code



Engineering
code



Extensions



Data

<https://github.com/PyTorchLightning/pytorch-lightning>

PYTORCH

```
# models
encoder = nn.Sequential(nn.Linear(28 * 28, 64), nn.ReLU(), nn.Linear(64, 3))
decoder = nn.Sequential(nn.Linear(3, 64), nn.ReLU(), nn.Linear(64, 28 * 28))

encoder.cuda(0)
decoder.cuda(0)

# download on rank 0 only
if global_rank == 0:
    mnist_train = MNIST(os.getcwd(), train=True, download=True)

# split dataset
transform=transforms.Compose([transforms.ToTensor(),
                              transforms.Normalize(0.5, 0.5)])
mnist_train = MNIST(os.getcwd(), train=True, download=True, transform=transform)

# train (55,000 images), val split (5,000 images)
mnist_train, mnist_val = random_split(mnist_train, [55000, 5000])

# The dataloaders handle shuffling, batching, etc...
mnist_train = DataLoader(mnist_train, batch_size=64)
mnist_val = DataLoader(mnist_val, batch_size=64)

# optimizer
params = [encoder.parameters(), decoder.parameters()]
optimizer = torch.optim.Adam(params, lr=1e-3)

# TRAIN LOOP
model.train()
num_epochs = 1
for epoch in range(num_epochs):
    for train_batch in mnist_train:
        x, y = train_batch
        x = x.cuda(0)
        x = x.view(x.size(0), -1)
        z = encoder(x)
        x_hat = decoder(z)
        loss = F.mse_loss(x_hat, x)
        print('train loss: ', loss.item())

        loss.backward()
        optimizer.step()
        optimizer.zero_grad()

# EVAL LOOP
model.eval()
with torch.no_grad():
    val_loss = []
    for val_batch in mnist_val:
        x, y = val_batch
        x = x.cuda(0)
        x = x.view(x.size(0), -1)
        z = encoder(x)
        x_hat = decoder(z)
        loss = F.mse_loss(x_hat, x)
        val_loss.append(loss)
    val_loss = torch.mean(torch.tensor(val_loss))
model.train()
```

PYTORCH LIGHTNING

Turn PyTorch into Lightning

Lightning is just plain PyTorch.

<https://github.com/PyTorchLightning/pytorch-lightning>



Ecosystem

- ▶ PyTorch lightning
- ▶ PyTorch geometric
- ▶ Hydra
- ▶ Horovod
- ▶ Skorch
- ▶ Captum
- ▶ And many others, see <https://pytorch.org/ecosystem/>


Moar stuff

- ▶ <https://pytorch.org/docs/stable/index.html>
- ▶ <https://pytorch.org/tutorials/beginner/ptcheat.html>
- ▶ <http://neuralnetworksanddeeplearning.com/chap2.html>
- ▶ <https://www.khanacademy.org/math/differential-calculus/dc-chain>
- ▶ <https://blog.paperspace.com/pytorch-101-understanding-graphs-and-automatic-differentiation/>

Conclusion

- ▶ PyTorch is a solid, flexible, production-ready foundation for real-life deep-learning applications
- ▶ Building blocks:
 - Tensors
 - Functions
- ▶ Dynamic graph automatic differentiation
 - CPU, GPU, TPU
- ▶ Rich ecosystem

Thank you!



Andrey Ustyuzhanin