

Nikita Kazeev



Generative Adversarial Network (GAN)

2021



Yandex



EPFL



Generative models in general

▶ Want:

- A neural network generator. When applied to noise z , samples the target distribution: $G(z) \sim P(x)$
- OR a neural network approximation of probability density
 - Not covered now, Artem Ryzhikov will teach them in a few sections

▶ Training:

- Minimize some divergence $D(P_{\text{train}} || Q_{\text{generated}})$

▶ Challenges:

- Both the training and generated distributions are not known explicitly



Consider an f-divergence

Can be computed directly Can be estimated by a classifier

$$D_f(P||Q) = \underbrace{\int dx \cdot q(x)}_{\text{Can be estimated by sampling}} \cdot f\left(\underbrace{\frac{p(x)}{q(x)}}_{\text{Can be estimated by a classifier}}\right)$$



Spherical generative adversarial network in vacuum

Repeat until convergence:

- ▶ Fit a classifier to estimate $\frac{p(x)}{q(x)}$
- ▶ Compute the value of an f-divergence between the real and generated data via Monte-Carlo sampling
- ▶ Train the generator: take a gradient descent step to minimize the distance

Challenges:

- ▶ Finite training data
- ▶ Finite computing resources



Divergence selection

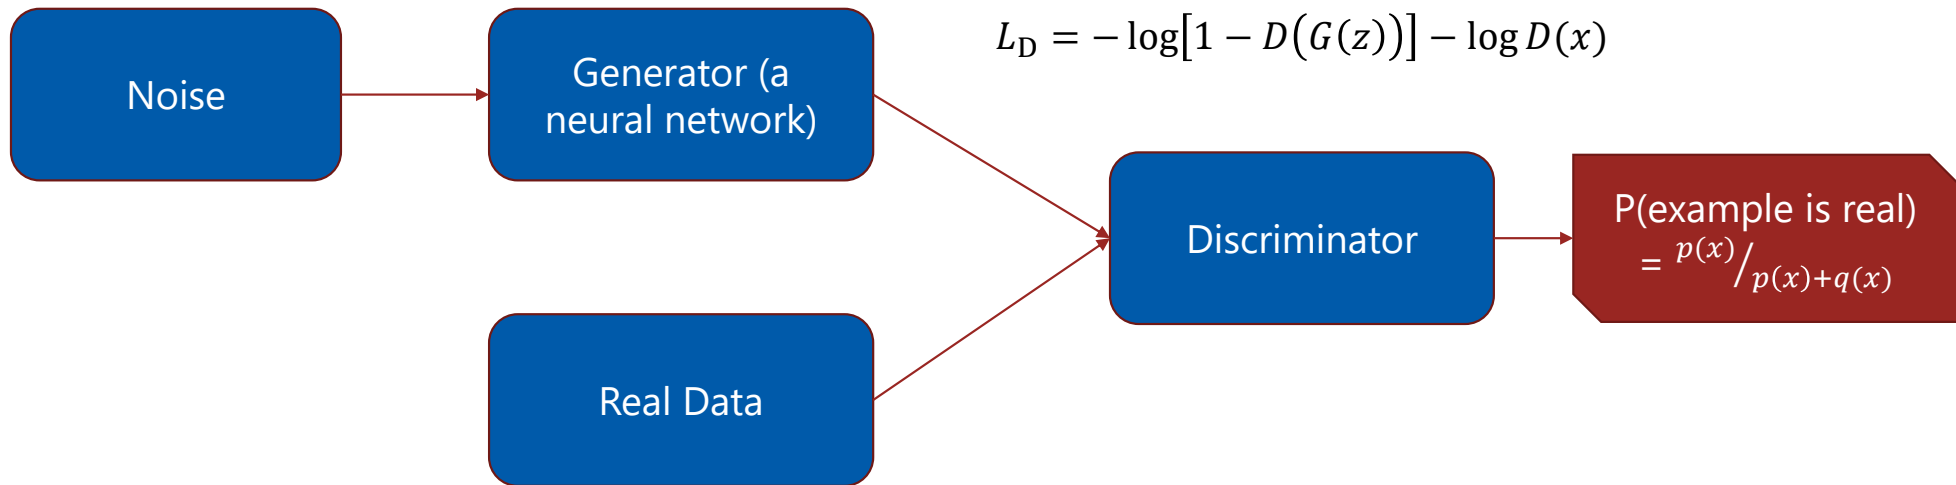
- ▣ Not KL, it is infinite if $\exists x: p(x) \neq 0, q(x) = 0$
 - For a complex distribution (e. g. images), it is almost certain to happen
- ▣ Most of the papers on GANs with f-divergencies use JS
 - It works well enough
 - No tangible benefit from using other divergencies
 - See a review paper: <https://arxiv.org/pdf/1606.00709.pdf>



JS GAN scheme

$$L_G = -\log D(G(z))$$

$$L_D = -\log[1 - D(G(z))] - \log D(x)$$



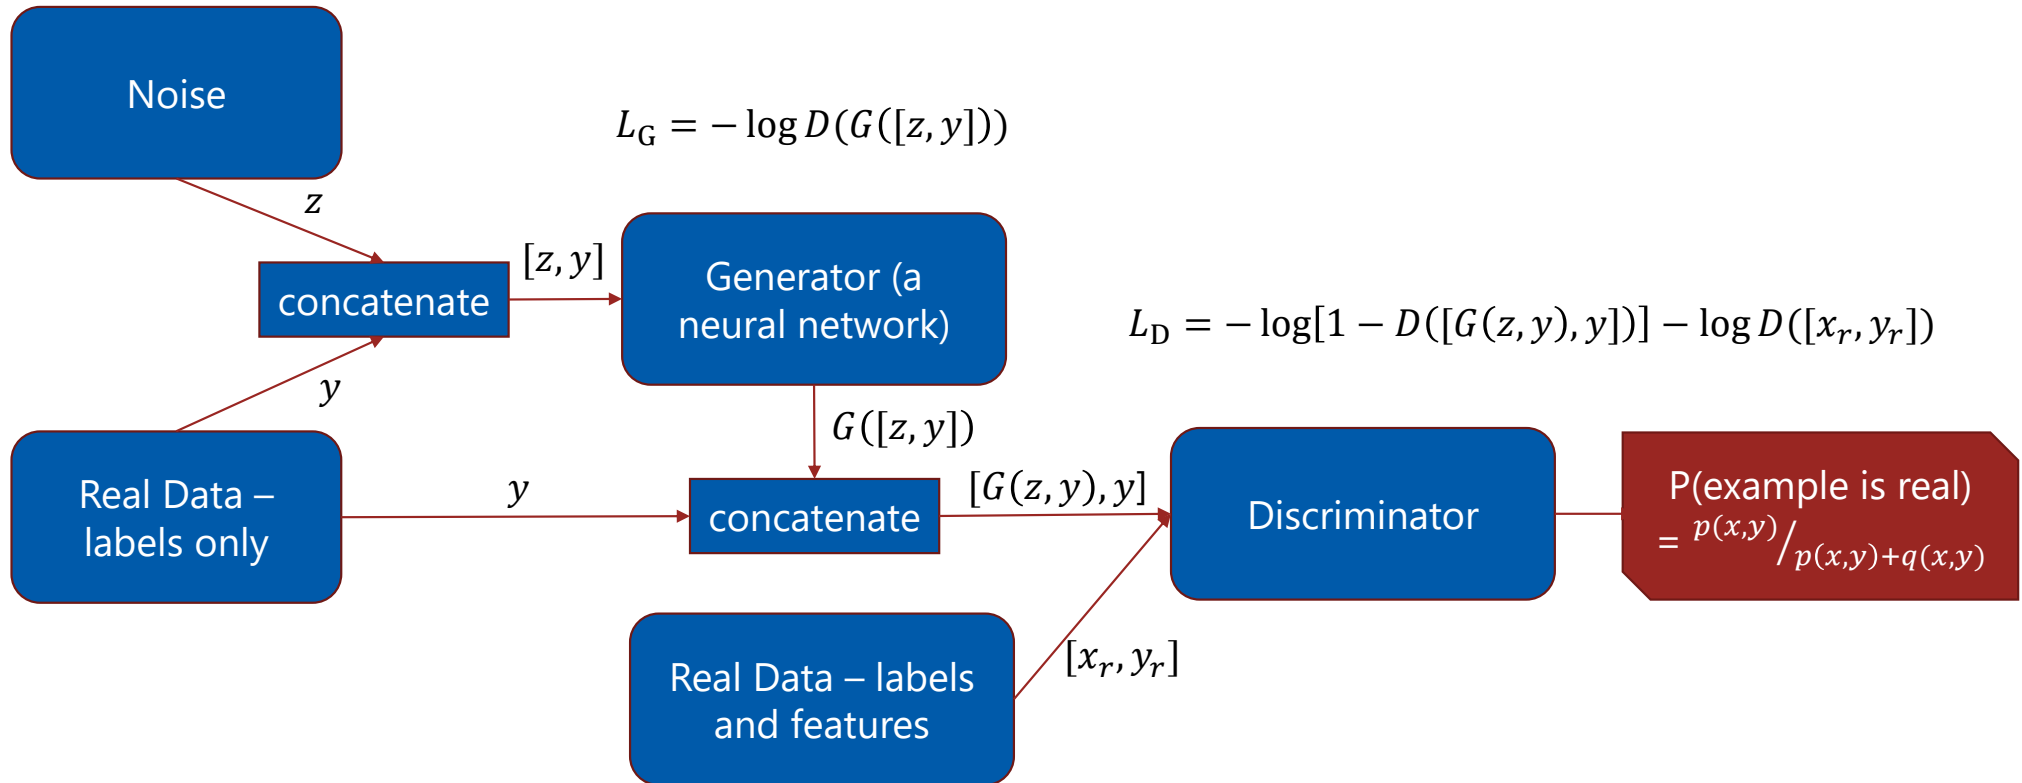
Training algorithm

Repeat while not good enough

- ▶ Train discriminator for N_d iterations:
 - sample a batch of training data x_r
 - sample a batch of noise z
 - compute discriminator loss $L_D = -\log[1 - D(G(z))] - \log D(x_r)$
 - take an optimization step for discriminator, minimizing L_D
- ▶ Train generator
 - sample a batch of noise z
 - compute generator loss $L_G = -\log D(G(z))$
 - take an optimization step for generator, minimizing L_G




Conditional GAN scheme



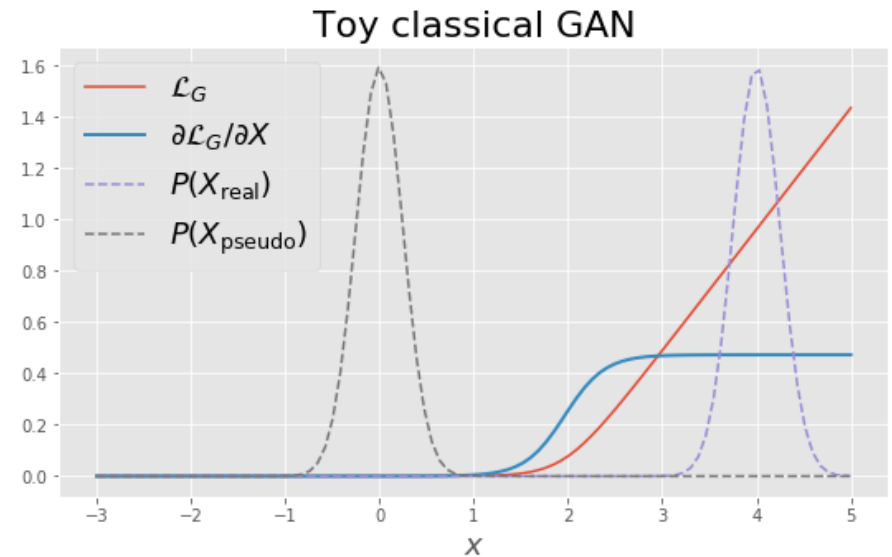
Problem #1: vanishing gradients

- ▶ Consider the case of disjoint support of real and generated data
- ▶ An ideal discriminator can perfectly tell the real and generated data apart:
$$D(G(z)) \approx 0$$



Problem #1: vanishing gradients

- ▣ $L_G = -\log D(G(z))$
- ▣ $dD(x)/dx \approx 0$ for generated x
- ▣ $dL_G(x)/dx \approx 0$ for generated x
- ▣ Generator can't train!



Fight for the gradients

Start with heavily restricted discriminator:

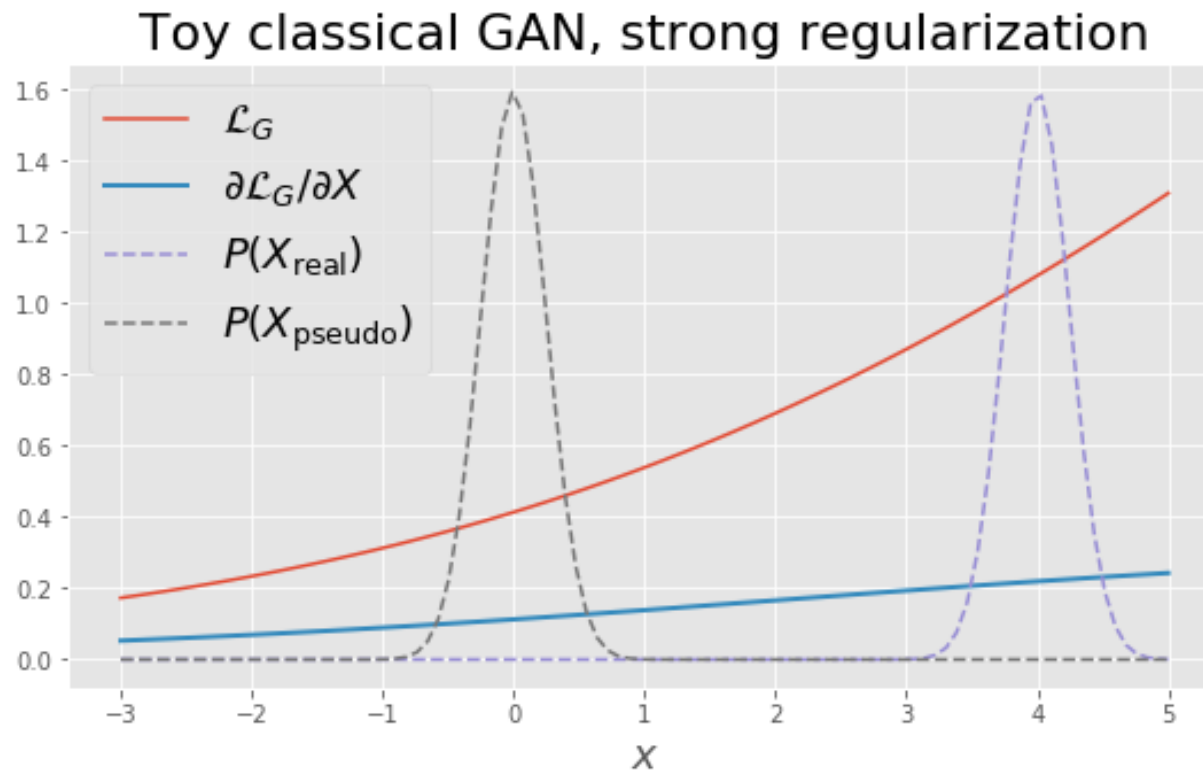
- ▶ don't train discriminator fully
- ▶ add noise to the samples:
 - nicely works for target on low-dimensional manifolds;
 - easy to control.
- ▶ discriminator regularization:
 - might interfere with the convergence.

As learning progresses gradually relax restrictions.


(or use a different class of divergencies – see the next lecture)



Fight for the gradients: regularization



Fight for the gradients: noise




Problem #2: mode collapse

- ▶ Generator only trains on the objects it generates:

$$L_G = -\log D(G(z))$$

- ▶ If there are different modes in the data distribution, the generator can get stuck in a local minimum of the discriminator
- ▶ Example. If a generator has learned to generate only "1", but perfectly, it might fail to learn to generate "5"



[Image: Adiga, Sudarshan, et al. "On the tradeoff between mode collapse and sample quality in generative adversarial networks." 2018 IEEE Global Conference on Signal and Information Processing \(GlobalSIP\). IEEE, 2018.](#)



Summary

- ▣ Classic generative adversarial networks with JS loss
 - Powerful method for generative modelling
 - Hard to train
 - Need to balance discriminator and generator power
 - Vanishing gradients
 - Mode collapse
 - Are the starting point of many image-specific advances
- ▣ See the next lecture for a newer approach based on integral probability measures



Thank you!



nkazeev@hse.ru



[kazeevn](https://www.telegram.me/kazeevn)



[hse_lambda](https://www.instagram.com/hse_lambda)

Nikita Kazeev

Acknowledgements

The presentation is based on the previous MLHEP editions: vanishing gradients slides by Maxim Borisyak